Tutorial Google Cloud Platform

Arthur Pesah, Sebastian Bujwid, Yonk Shi

March 25, 2019

Contents

1 Introduction

The main option to have access to a GPU for your deep learning project is to use Google Cloud Platform (GCP). We've obtained some GCP educational grant for this course, each student's email is entitled $50 for the course (If your team has three members, you will have a combined $150 credits - which can correspond to around 360 hours of using NVIDIA K80!). To claim your credit, simply fill out this request form with your KTH email address (since we have a limited number of coupons, please do not redistribute this link!). Follow the instructions in the email, once you have setup the coupon, you will see

a billing account named **DD2424**. Make sure to use this DD2424 billing account when you create a new project / applying billing to an existing project.

Google also allows everyone to get a free trial account with $300 of credits. If you run out of credits with us, don't hesitate to start using this trial.

2 How to add new members to your project

You might want other members of your team to be able to access to GCP. In that case, it's pretty straight-forward to add new members. Go to *IAM & admin / IAM*:

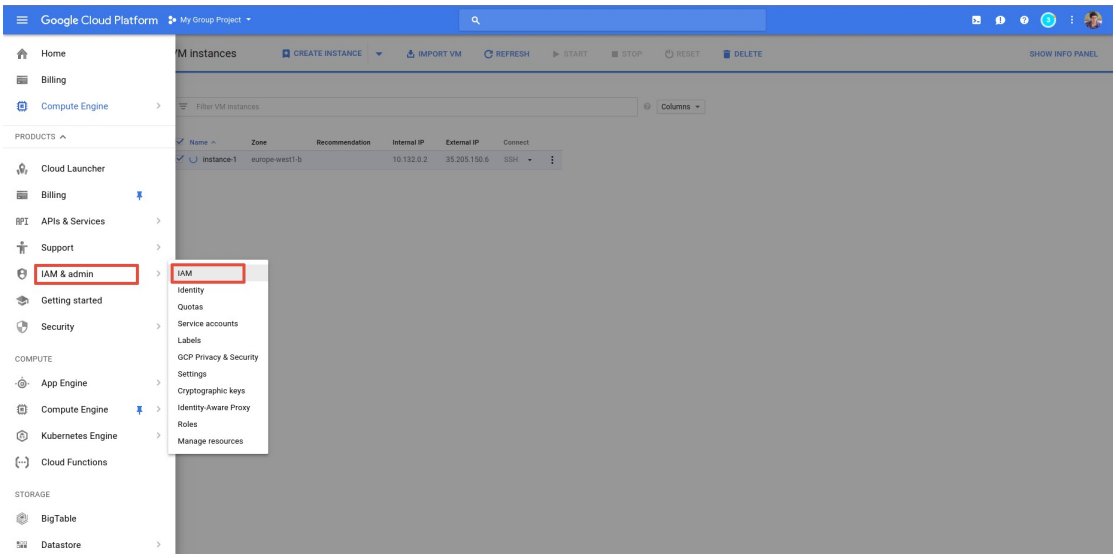

Then click on *Add*:

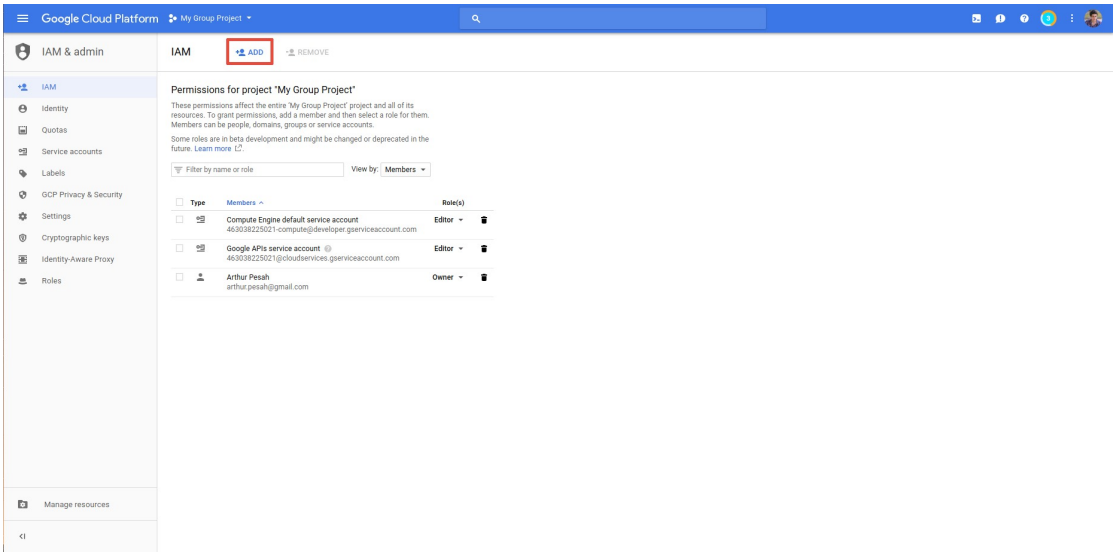

Add the google address of your other team members and select *Project/Owner* as a role:

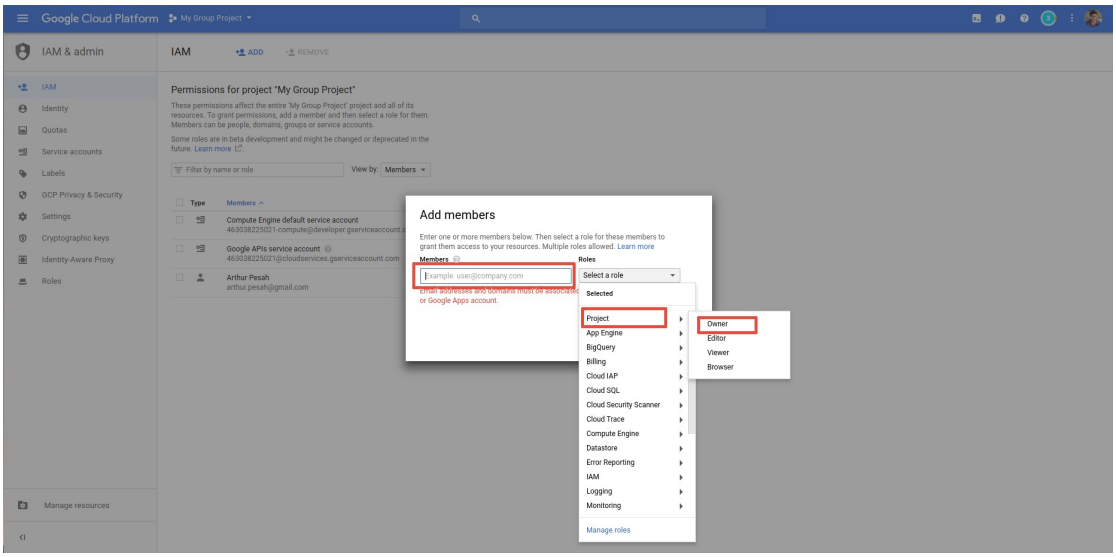

3 GPU Quota & Setup

If you are planning on using GPU resources, you will have to set GPU quota for your project. It is likely that the GPU quota for your project is 0 by default. You can check that in *IAM & admin* → *Quotas*, under service called *Compute Engine API - GPUs (all regions)*. If it says 0 you need to increase it to be able to deploy GPU machines. You can do that by selecting the service and clicking *EDIT QUOTAS*.

When creating the virtual machine/choosing quota, we recommend **K80 GPU** over expensive GPUs (V100, P100 etc). K80s are a fraction of the cost comparing to other GPUs, and their performances are comparable. In our experience, two K80s will outperform a single V100, yet still cheaper.

Note: do not apply for higher quota than you really need. The higher the quota you apply for, the more time it can take for Google to process the request. Requesting the limit of 1 should be approved very quickly.

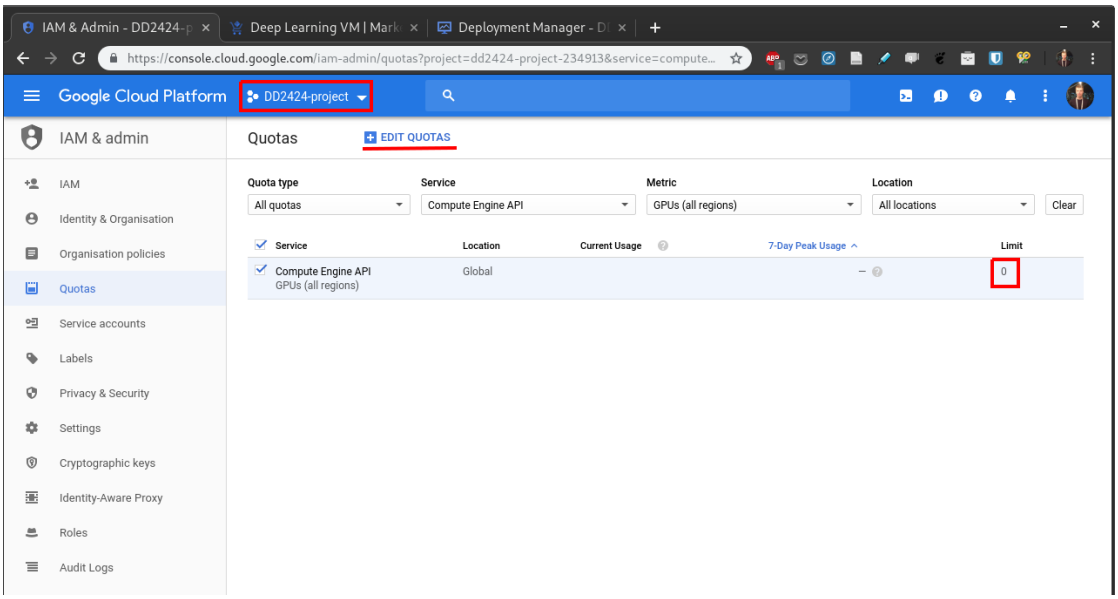

4 How to create a Deep Learning virtual machine

The most convenient way to do use GCP resources for Deep Learning is to use existing Deep Learning image templates. They are easy and quick to setup and come with software packages such as Tensorflow or PyTorch and have pre-installed NVIDIA drivers.

To deploy a new *Deep Learning VM* go to the GCP Console page and search for *Deep Learning VM* in the search bar at the top of the page.

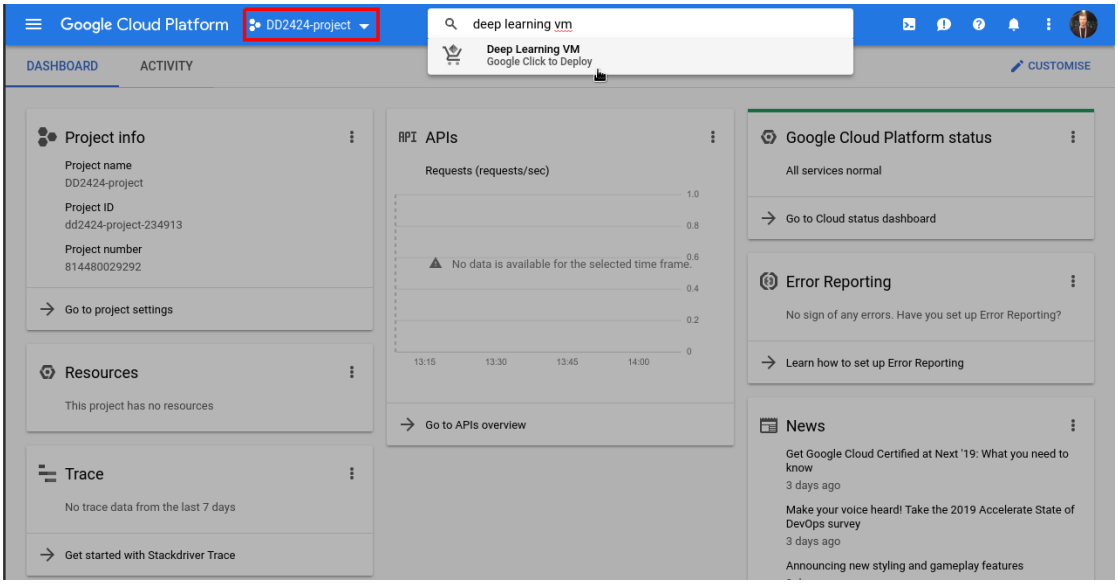

When selecting the image **make sure that you have the right project active** - you can change the project next to GCP logo on top left corner.

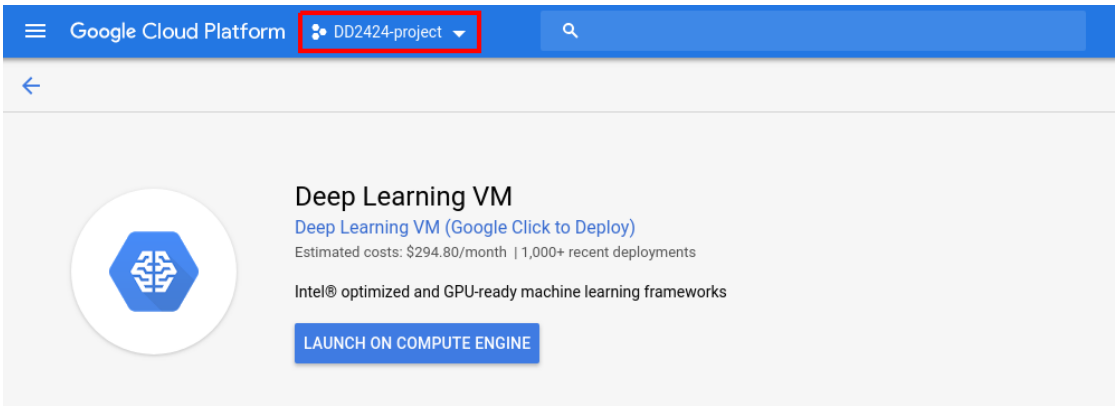

After clicking *Launch on compute engine* you will be moved to the settings of the VM. You can choose the size of the machine and the disk space according to your needs (keep in mind that it affects the price a little bit). If you're using large datasets that require a parallel reader and fast pre-fetching, you might also want to choose a better CPU (with typically more cores). Select the GPU you need (we advise to use K80) and the framework you intend to use, e.g *Intel optimized TensorFlow* or *PyTorch*. **Make sure that you select the option to install NVIDIA GPU driver**.

Optionally, use can change the zone to *europe-west1-b* or *europe-west1-d* - since the servers are closer the connection should have smaller delays.

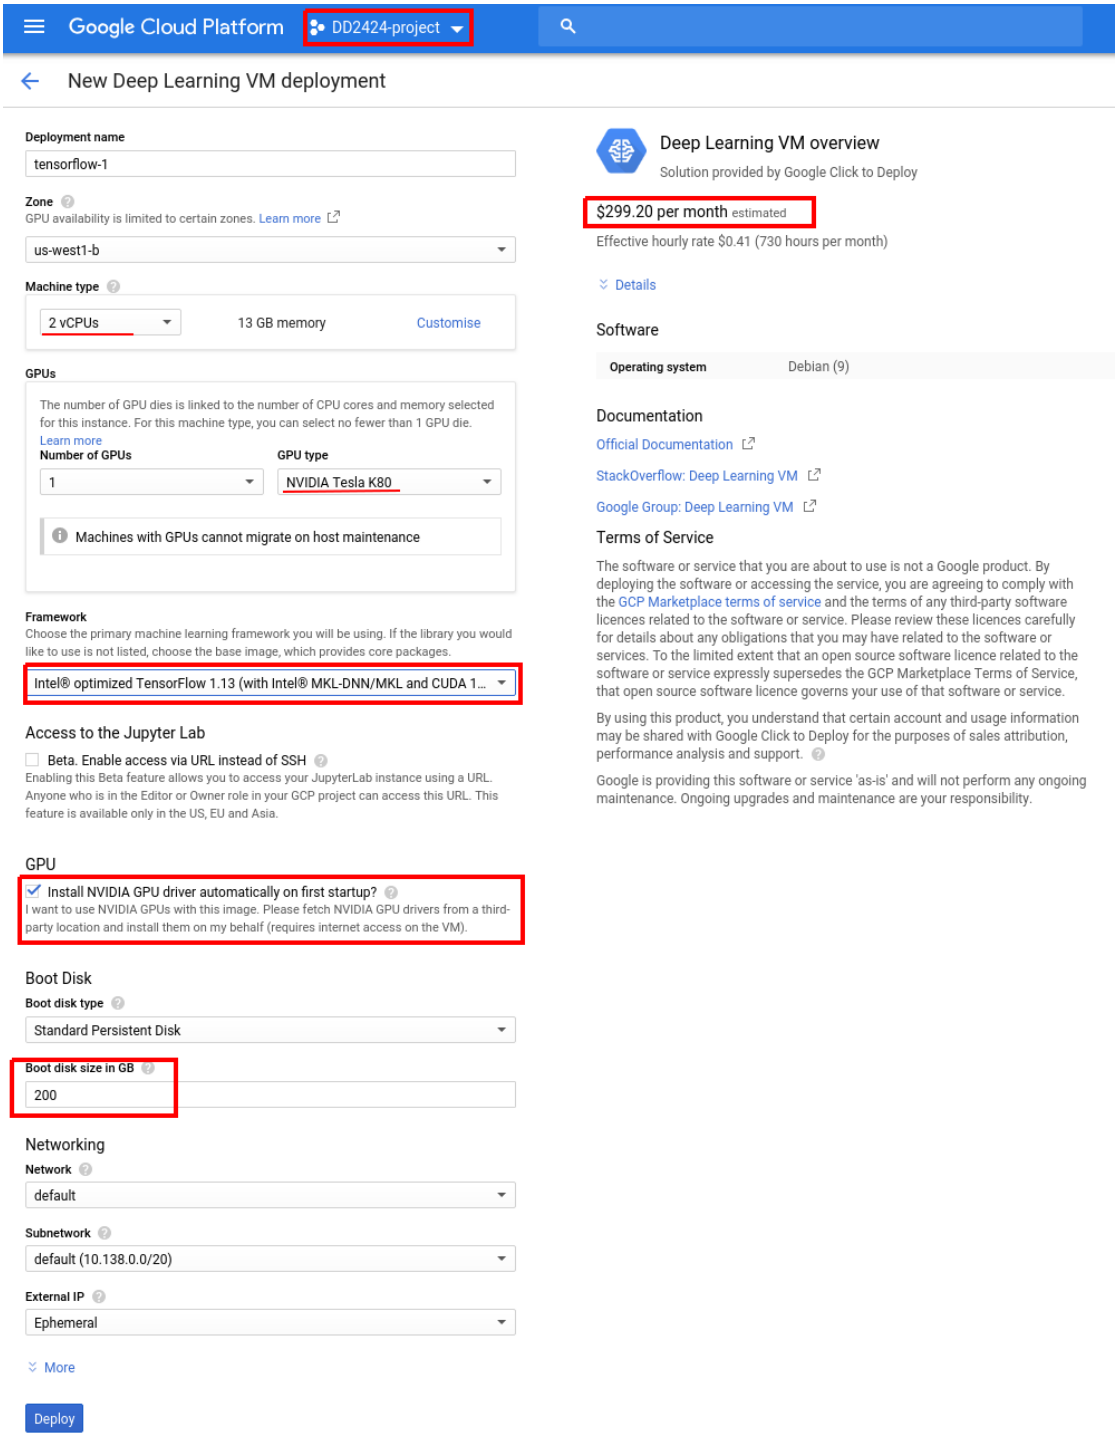

After you click *Deploy*, each hour is counted on your budget, so don't forget to **STOP THE VM WHEN YOU'RE NOT USING IT** (instructions at the end).

The next step to be able to connect is to handle your ssh keys. The instructions depend on your OS and are well explained in the GCP documentation. connect is to handle your ssh keys. The instructions depend on your OS and are well explained on the GCP documentation. You can create as many ssh keys as you want, so don't hesitate to repeat the instructions for everyone in your group connecting from his own computer. Here is a summary for Linux users:

```
$ ssh-keygen -t rsa -f ~/.ssh/gcp-key -C [GOOGLE EMAIL ADDRESS]
```

You can use the default parameters for the creation of the key. Then:

```
$ cat ~/.ssh/gcp-key.pub
```

Copy the entire output (including your email address that should appear at the end of the key) and go back to GCP. Click on *Compute Engine / Metadata*:

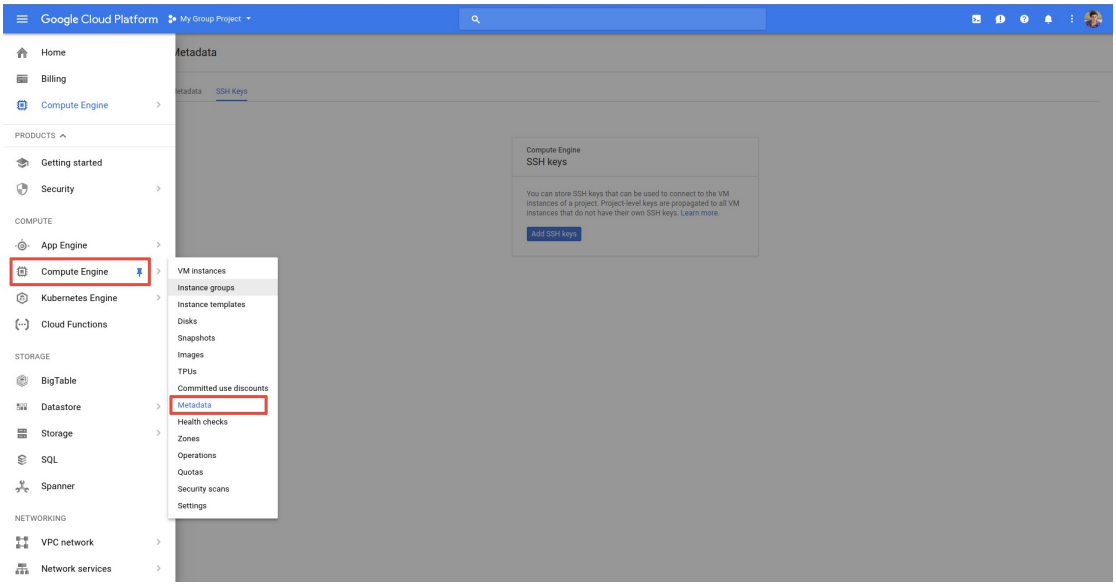

Then click on *SSH Keys / Add SSH keys*:

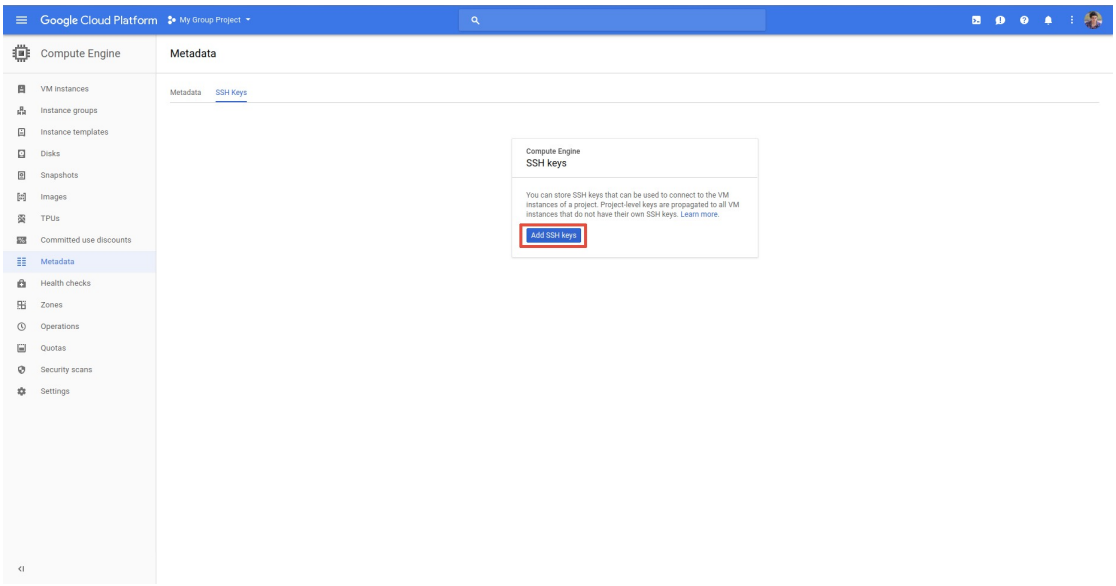

And paste your key in the form. Delete any extra space or newline, the key should end at the end of the email address. Then save:

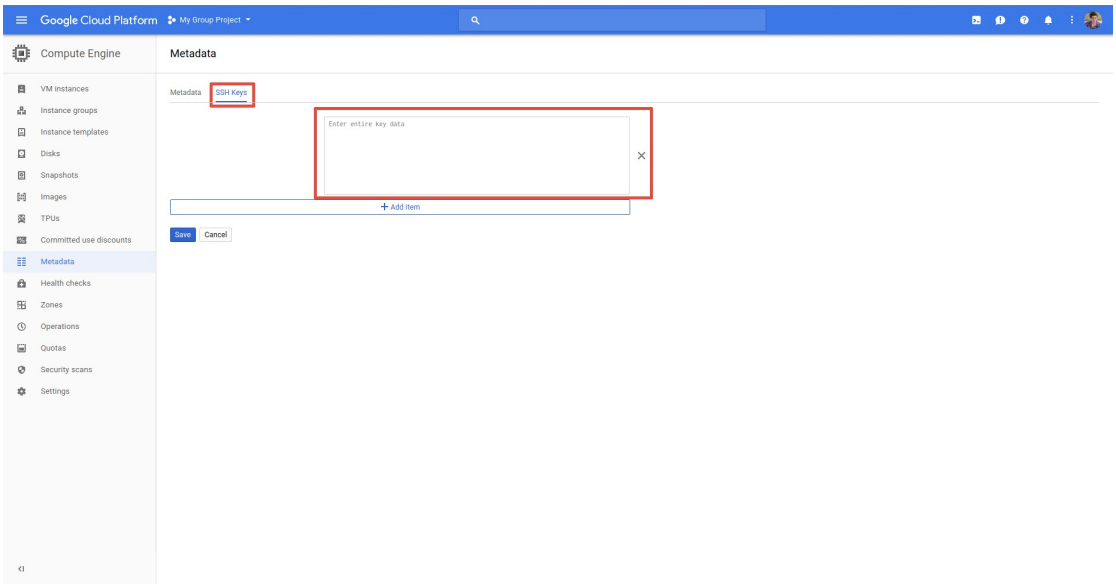

You're now able to connect to your VM. One way to do it is to use the SSH button on the right (it will open a new window with a web terminal). But if you prefer connecting from your own terminal, just copy the external IP address like below:

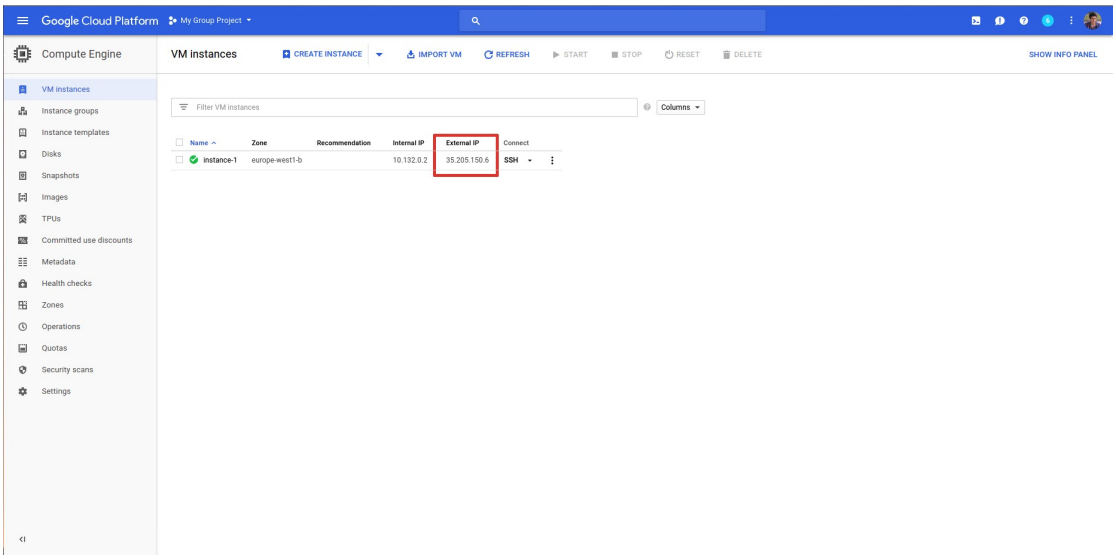

And connect to your machine using a command like that:

`$ ssh -i ~/.ssh/gcp-key [google username]@[external ip address]`

You can also copy your project to the machine using scp:

`$ scp -i ~/.ssh/gcp-key -r [projectFolder] [google username]@[external ip address]:~`

To stop the machine when you are not using it, just select it and click on STOP:

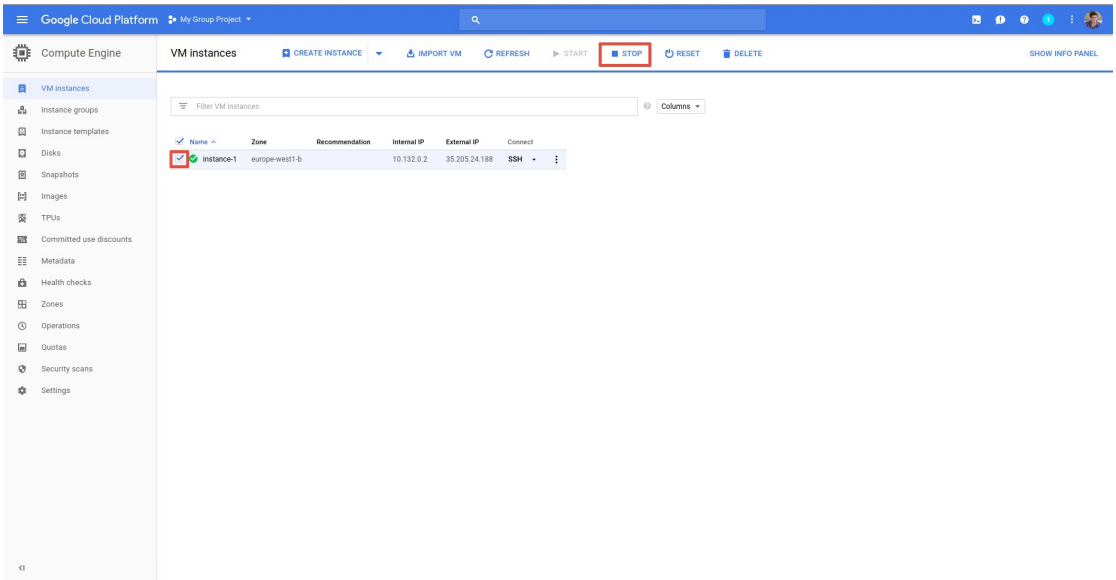

5 Optional - installing specific versions of packages

5.1 Using conda

In case you need very specific versions of Python deep learning packages the easiest way to get them is through Miniconda. *conda* installs necessary dependencies for you (such as cudnn, cuda-toolkit, etc.).

Note: Packages installed with *conda* can be faster than packages installed with *pip*, because of more compilation flags enabled (support for Intel MKL-DNN, etc.).

If the command *conda* is not available you can install it by executing:

```
$ wget https://repo.anaconda.com/miniconda/Miniconda3-latest-
    Linux-x86_64.sh
$ bash Miniconda3-latest-Linux-x86_64.sh}
```

During the installation process confirm to initialize Miniconda:

```
Do you wish the installer to initialize Miniconda3
in your /home/sebastian/.bashrc ? [yes|no]
[no] >>> yes
```

To be able to access *conda* you need to log out and log in to the machine again through ssh. After that you should be able to use *conda* to create new environments:

```
# e.g. installing TensorFlow 1.10.0
$ conda create -n myenv1 tensorflow-gpu=1.10.0
$ conda activate myenv1
$ python -c "import tensorflow as tf; print(tf.__version__)"
  1.10.0

# e.g. installing Pytorch 0.3.1
$ conda create -n myenv2 pytorch=0.3.1
$ conda activate myenv2
$ python -c "import torch; print(torch.__version__)"
  0.3.1

# e.g. installing Caffe
$ conda create -n myenv3
$ conda activate myenv3
$ conda install caffe-gpu
$ caffe --version
caffe version 1.0.0
```

```
Debug build (NDEBUG not #defined)
```

To install more packages you can use *conda install* or *pip install*.

5.2 Using GCP Marketplace image templates

An alternative approach for obtaining more specific software is to use pre-defined image templates from GCP Marketplace. Use the search bar to find software packages that you are interested in. For example searching for *Caffe* you can find image *Caffe Python 3.6 NVidia GPU Production* by *Jetware*. If you find an image template that meets your particular requirements click *Launch on compute engine* and follow relevant instructions from Section 4.

Note: Using some image templates might not be free, make sure to check the pricing when launching.